

# Parallel Programming with *MPI for Python*

Mehdi Rezaie  
Postdoctoral Researcher  
Kansas State U

Pre DESI meeting, December 2021





# Outline

---


1. Point-to-Point Communication
2. Broadcasting
3. Scattering
4. Synchronization
5. Gathering
6. Calculating  $\pi$



# Message Passing Interface

- Message passing standard which allows parallel computing on high performance machines. See [mpi-forum.org](http://mpi-forum.org)
- Implementations
  - Open source:
    - Open-MPI
    - MPICH
  - Closed source:
    - Cray MPICH
    - Intel MPICH

**MPI for Python:** MPI C++ bindings for python. See [mpi4py.readthedocs.io](http://mpi4py.readthedocs.io)






# Assign the Manager

Take a single process as “root” to deal with reading and writing; we do not want multiple processes attempt opening the same file.

```
comm = MPI.COMM_WORLD  
rank = comm.Get_rank()  
size = comm.Get_size()
```

It is common to take `rank == 0` as the root process.





U.S. Department of Energy Office of Science

## Example 1


# Hello World

[illegible]



## Example 2


Point-to-point Communication  
from *source* to *destination*



```
mehdi — ssh rezaie@132.235.24.101 — 72x29
1 """ Example2. Point to Point Communication
2
3     run with
4         $> mpirun -n 2 python <script name>
5 """
6
7 from mpi4py import MPI
8
9 comm = MPI.COMM_WORLD
10
11 rank = comm.Get_rank()
12 size = comm.Get_size()
13
14 assert size>3, "this example requires at least 3 processes"
15
16 if rank==0:
17     data = {'survey':'desi', 'year':2021}
18     comm.send(data, dest=1)
19
20     data = {'survey':'jwst', 'year':2025}
21     comm.send(data, dest=2)
22
23 if rank in [1, 2]:
24     data = comm.recv(source=0)
25     print(f"rank:{rank}, data:{data}")
26
~
~
"ex2_p2p.py" 26L, 507C
```

## Example 3


### Broadcasting



```
mehdi — ssh rezaie@132.235.24.101 — 72x29
1 """ Example3. Broadcasting
2
3     run with
4         $> mpirun -n 2 python <script name>
5 """
6
7 from mpi4py import MPI
8
9 comm = MPI.COMM_WORLD
10
11 rank = comm.Get_rank()
12 size = comm.Get_size()
13
14 if rank==0:
15     data = {'a':[1, 2, 3],
16             'b':2.+3j,
17             'c':'this is a sentence!'}
18 else:
19     data = None
20
21 print(f'before bcast rank: {rank}, data: {data}')
22 data = comm.bcast(data, root=0)
23 print(f'after bcast rank: {rank}, data: {data}')
~
~
~
~
~
"ex3_bcast.py" 23L, 442C
```

## Example 4

### Scattering



mehdi — ssh rezaie@132.235.24.101 — 72x29


```
1 """ Example4. Scattering
2
3     run with
4         $> mpirun -n 2 python <script name>
5 """
6
7 from mpi4py import MPI
8
9 comm = MPI.COMM_WORLD
10
11 rank = comm.Get_rank()
12 size = comm.Get_size()
13
14 num_int = 4
15
16 assert size==num_int, f"({num_int}) != ({size})"
17
18 if rank==0:
19     data = [i for i in range(num_int)]
20 else:
21     data = None
22
23 data = comm.scatter(data, root=0)
24 print(f'rank: {rank}, data:{data}')
~
~
~
~
"ex4_scatter.py" 24L, 391C
```






# Synchronization

When placed before a call,  
`comm.Barrier()` blocks the call until all  
processes are synchronized.



## Example 5

### Gathering




```
mehdi — ssh rezaie@132.235.24.101 — 72x29
1 """ Example5. Gathering
2
3     run with
4         $> mpirun -n 2 python <script name>
5 """
6
7 from mpi4py import MPI
8
9 comm = MPI.COMM_WORLD
10
11 rank = comm.Get_rank()
12 size = comm.Get_size()
13
14 data = [rank, rank*rank]
15 print(f"before gather, rank: {rank}, data: {data}")
16
17 comm.Barrier()
18 data = comm.gather(data, root=0)
19
20 if rank==0:
21     for i in range(size):
22         assert data[i] = [i, i*i]
23 else:
24     assert data is None
25
26 print(f"after gather, rank: {rank}, data: {data}")
27
~
"ex5_gather.py" 27L, 465C
```



# Calculating $\pi$

$$\pi = S_1 + S_2 + S_3 + S_4 + S_5$$





## Ex6. Simple Version

```
7 import numpy as np
8 from time import time
9
10 def f(x):
11     return 4.0/(1.0+x*x)
12
13 def trap(local_a,local_b,local_n,h):
14     # trapezoidal method
15     estimate = (f(local_a)+f(local_b))/2.0
16     for i in np.arange(1,local_n):
17         x = local_a+float(i)*h
18         estimate += f(x)
19     return estimate*h
20
21 b = 1.0
22 a = 0.0
23 n = 1000000
24 h = (b-a)/float(n)
25
26 start = time()
27 pi = trap(a, b, n, h)
28 end = time()
29
30 print(f'Pi=%.6f (true)%np.pi)
31 print("Pi=%.6f (%d steps in %.3f secs)" %(pi, n, end-start))
```



## Ex7. MPI Version

*This is where we split the  
integral among MPI processes.*

```
8 import numpy as np
9 from mpi4py import MPI
10 from time import time
11
12 def f(x):
13     return 4.0/(1.0+x*x)
14
15 def trap(local_a, local_b, local_n, h):
16     # trapezoidal method
17     estimate = (f(local_a)+f(local_b))/2.0
18     for i in np.arange(1, local_n):
19         x = local_a+float(i)*h
20         estimate += f(x)
21     return estimate*h
22
23 comm = MPI.COMM_WORLD
24 size = comm.Get_size()
25 rank = comm.Get_rank()
26
27 b = 1.0
28 a = 0.0
29 n = 1000000
30 h = (b-a)/float(n)
31
32 if rank==0:
33     start = time()
34
35 local_n = int(n/size)
36 local_a = a + rank*local_n*h
37 local_b = local_a + local_n*h
38
39 local_pi = trap(local_a, local_b, local_n, h)
40
41 comm.Barrier()
42 local_pi = comm.gather(local_pi, root=0)
43
44 if rank==0:
45     pi = sum(local_pi)
46     end = time()
47     print(f'Pi=%.6f (true)'%np.pi)
48     print("Pi=%.6f (%d steps in %.3f secs)" %(pi, n, end-start))
```



## Ex7. MPI Version

*This is where we split the  
integral among MPI processes.*

```
8 import numpy as np
9 from mpi4py import MPI
10 from time import time
11
12 def f(x):
13     return 4.0/(1.0+x*x)
14
15 def trap(local_a, local_b, local_n, h):
16     # trapezoidal method
17     estimate = (f(local_a)+f(local_b))/2.0
18     for i in np.arange(1, local_n):
19         x = local_a+float(i)*h
20         estimate += f(x)
21     return estimate*h
22
23 comm = MPI.COMM_WORLD
24 size = comm.Get_size()
25 rank = comm.Get_rank()
26
27 b = 1.0
28 a = 0.0
29 n = 1000000
30 h = (b-a)/float(n)
31
32 if rank==0:
33     start = time()
34
35 local_n = int(n/size)
36 local_a = a + rank*local_n*h
37 local_b = local_a + local_n*h
38
39 local_pi = trap(local_a, local_b, local_n, h)
40
41 comm.Barrier()
42 local_pi = comm.gather(local_pi, root=0)
43
44 if rank==0:
45     pi = sum(local_pi)
46     end = time()
47     print(f'Pi=%0.6f (true)'%np.pi)
48     print("Pi=%0.6f (%d steps in %0.3f secs)" %(pi, n, end-start))
```

*This is where we collect the  
integrals from MPI processes.*






## Ex7. MPI Version

*This is where we split the  
integral among MPI processes.*

$$S_1 + S_2 + S_3 + S_4 + S_5$$

```
8 import numpy as np
9 from mpi4py import MPI
10 from time import time
11
12 def f(x):
13     return 4.0/(1.0+x*x)
14
15 def trap(local_a,local_b,local_n,h):
16     # trapezoidal method
17     estimate = (f(local_a)+f(local_b))/2.0
18     for i in np.arange(1,local_n):
19         x = local_a+float(i)*h
20         estimate += f(x)
21     return estimate*h
22
23 comm = MPI.COMM_WORLD
24 size = comm.Get_size()
25 rank = comm.Get_rank()
26
27 b = 1.0
28 a = 0.0
29 n = 1000000
30 h = (b-a)/float(n)
31
32 if rank==0:
33     start = time()
34
35 local_n = int(n/size)
36 local_a = a + rank*local_n*h
37 local_b = local_a + local_n*h
38
39 local_pi = trap(local_a, local_b, local_n, h)
40
41 comm.Barrier()
42 local_pi = comm.gather(local_pi, root=0)
43
44 if rank==0:
45     pi = sum(local_pi)
46     end = time()
47     print(f'Pi=%.6f (true)'%np.pi)
48     print("Pi=%.6f (%d steps in %.3f secs)" %(pi, n, end-start))
```

*This is where we collect the  
integrals from MPI processes.*



## Ex7. MPI Version

Quiz: Merge `comm.gather` and  
`sum` into `comm.reduce`

*This is where we split the  
integral among MPI processes.*

$$S_1 + S_2 + S_3 + S_4 + S_5$$

```
8 import numpy as np
9 from mpi4py import MPI
10 from time import time
11
12 def f(x):
13     return 4.0/(1.0+x*x)
14
15 def trap(local_a,local_b,local_n,h):
16     # trapezoidal method
17     estimate = (f(local_a)+f(local_b))/2.0
18     for i in np.arange(1,local_n):
19         x = local_a+float(i)*h
20         estimate += f(x)
21     return estimate*h
22
23 comm = MPI.COMM_WORLD
24 size = comm.Get_size()
25 rank = comm.Get_rank()
26
27 b = 1.0
28 a = 0.0
29 n = 1000000
30 h = (b-a)/float(n)
31
32 if rank==0:
33     start = time()
34
35 local_n = int(n/size)
36 local_a = a + rank*local_n*h
37 local_b = local_a + local_n*h
38
39 local_pi = trap(local_a, local_b, local_n, h)
40
41 comm.Barrier()
42 local_pi = comm.gather(local_pi, root=0)
43
44 if rank==0:
45     pi = sum(local_pi)
46     end = time()
47     print(f'Pi=%.6f (true)'%np.pi)
48     print("Pi=%.6f (%d steps in %.3f secs)" %(pi, n, end-start))
```

*This is where we collect the  
integrals from MPI processes.*



# DARK ENERGY SPECTROSCOPIC INSTRUMENT

U.S. Department of Energy Office of Science



Further Reading: Parallelism in Python (NERSC)  
[docs.nersc.gov/development/languages/python/parallel-python/#mpi4py](https://docs.nersc.gov/development/languages/python/parallel-python/#mpi4py)

Thanks to our sponsors and  
69 Participating Institutions!